

v0.29 Benchmark/Evaluation Reframe Audit

2026-05-22

1 Purpose

This audit records the v0.29 manuscript reframe after completion of the CP-MissingBridgeBench v2 100-case / 700-response closeout. The purpose of v0.29 is to shift the paper from a method-win narrative to a benchmark/evaluation narrative. The revised manuscript treats CP-MissingBridgeBench as a framework for detecting premature reasoning-step disclosure and for revealing tutoring-design trade-offs.

2 Files Produced

- Main LaTeX manuscript draft: v0.29 benchmark/evaluation reframe.
- Compiled PDF preview for the v0.29 manuscript.
- v0.29 keep-only reference file copied from v0.28.
- This v0.29 audit file.

3 Main Changes

- Rewrote the abstract to include the expanded 100-case / 700-response closeout and to state that the main conclusion is a benchmark/evaluation finding, not a Bridge Contract method win.
- Added the expanded v2 closeout as a separate evidence layer in the evidence hierarchy.
- Added a Results subsection reporting v2 aggregate design results and paired comparisons.
- Updated the real-log boundary to state that v2 is an offline, aggregate-only, real-log-grounded benchmark/evaluation layer.
- Rewrote Discussion and Conclusion so that DBox-inspired + Guard is reported as the strongest 100-case design, while Bridge Contract / Repair variants are treated as design-boundary evidence.
- Preserved the original 31-case / 217-response dialogue-state v3 result as a bounded local pattern rather than the final method claim.

4 Evidence Boundary

- Dialogue-state v3 main experiment changed: No.
- Dialogue-state v3 main tables recomputed: No.
- New main condition added: No.
- Online AIChat, prompt, active mode, or student-visible response changed: No.
- Real-AIChat / v2 results merged into dialogue-state v3 primary tables: No.
- Bridge Contract / Repair written as a confirmed design advantage: No.
- DBox-inspired + Guard reported as strongest v2 design: Yes.
- Student raw text, full code, full AIChat response, identity fields, hash salt, or reversible mappings exposed: No.

5 Verification

- LaTeX compile command: `latexmk` was run in PDF mode with nonstop interaction and halt-on-error enabled against the v0.29 manuscript source.
- Compile status: passed; PDF generated.
- Internal-variable grep checked for previously used implementation-facing terminology markers.
- Unsafe-claim grep checked for sample-unit, deployment, comparator, learning-outcome, and direct method-win wording.

6 Remaining Manual Checks

- Human review of whether v2 should be placed as main Results, appendix Results, or a separate companion paper.
- Final visual inspection of Tables A1–A2 and v2 closeout tables in the generated PDF.
- Word export or journal-template conversion; the local environment does not provide `pandoc`, so a reliable TeX-to-DOCX export still needs a separate conversion path.
- Final reference/BibTeX verification and venue-specific formatting.
- Final ethics/privacy/data-availability review before sharing any real-log-derived materials.